

ArcWarden: a GPU-native particle-in-cell framework for high-throughput kinetic plasma simulations

Donglai Ma

Department of Atmospheric and Oceanic Sciences, UCLA
dma96@atmos.ucla.edu

Draft v0.2 — August 12, 2026

Abstract

We present ArcWarden, a GPU-native 2D3V particle-in-cell (PIC) framework designed for high-throughput kinetic plasma simulation on a single modern GPU. Rather than porting the established multi-node PIC design to accelerators, ArcWarden is built around what current GPUs provide: a last-level cache large enough to hold the entire 2D field state, fast shared-memory atomics, and inexpensive wide kernel launches. The particle step is a single fused kernel (field gather, Boris push, current scatter, and migration), currents accumulate in shared-memory tile aprons, and the particle sort that enables tiling is amortized over tens of steps and *stale-tolerant* — correctness never depends on sort recency. On identical physics (an anisotropy-driven whistler and electron-acoustic-wave problem with 156 million particles and 2.1×10^5 steps), ArcWarden is $2.6\times$ faster than the CUDA branch of the OSIRIS framework on the same device, sustaining 1.55×10^{10} particle-steps per second. A modular architecture lets electrostatic, spectral-Darwin, and full-Maxwell field solvers — together with inhomogeneous background fields, loss-cone and anisotropic distributions, δf and cold-fluid species, and absorbing boundaries — share one tested GPU particle engine, with every physics setup expressed as a runtime input deck. ArcWarden passes a three-level validation ladder: textbook verification (dispersion, growth rates, conservation), reproduction of the published nonlinear trapping simulations of An et al. [PRL 122, 045101 (2019)], and, as an advanced capability demonstration, self-consistent generation of coherent whistler-mode waves with rising-tone frequency chirping in an inhomogeneous magnetic field, following Tao et al. [GRL 44, 3441 (2017)].

Contents

1	Introduction	2
2	ArcWarden architecture	3
2.1	Overall design	3
2.2	Particle representation and memory layout	4
2.3	Fused gather–push–deposit kernel	4
2.4	Current deposition: shared-memory tile aprons	4
2.5	Particle sorting: amortized and stale-tolerant	5
2.6	Field solvers	5
2.7	Diagnostics and I/O	5
2.8	Correctness gates	5

3	Numerical models	6
3.1	Particle equations and interpolation	6
3.2	Electrostatic and Darwin spectral solvers	6
3.3	Full-Maxwell FDTD solver	6
3.4	One-dimensional electron-hybrid model	6
3.5	δf representation	6
4	Verification	7
4.1	Cross-branch validation	7
5	GPU performance	7
5.1	Head-to-head against OSIRIS-CUDA	9
5.2	Ablation: where the speedup comes from	9
5.3	The classic design is overhead-bound on this hardware	9
5.4	Kernel-level budget	11
5.5	Model choice as a performance multiplier: the Darwin branch	11
6	Modular physics capabilities	11
7	Advanced nonlinear demonstrations	12
7.1	Published nonlinear benchmark: whistler-driven electron trapping (An et al. 2019) .	12
7.2	Capability demonstration: self-consistent coherent whistler waves with frequency chirping in an inhomogeneous plasma	14
8	Discussion	14
9	Conclusions	16

1 Introduction

Kinetic simulation is the primary tool for collisionless plasma physics at electron scales: whistler-mode chorus and hiss, electron-cyclotron harmonic waves, electron-acoustic waves, and the associated time-domain structures all require a particle-in-cell (PIC) treatment [1, 2]. The questions are statistical — growth from shot noise, nonlinear saturation, resonant trapping — and demand hundreds of particles per cell over 10^5 – 10^6 time steps. Raw particle throughput (particle-steps per second) therefore decides which problems are reachable, and how many points of a parameter scan a researcher can afford.

Mature GPU-capable PIC frameworks exist — OSIRIS [3], VPIC [4], PIConGPU [5], WarpX [6], and the algorithm studies of Decyk and Singh [7]. They are engineered for a specific and important regime: very large problems decomposed across many nodes, with the GPU port grafted onto a design vocabulary formed by the hardware of the early 2010s — small caches, slow global atomics, and expensive kernel launches, answered by tiling the grid, binding particles to tiles in fixed-size chunks, and paying for that locality with per-step particle bookkeeping (find the movers, exchange, compact).

Meanwhile the hardware under a single workstation has changed character. A consumer NVIDIA Blackwell GPU (RTX 5090) offers 96 MB of L2 cache, 1.8 TB/s of memory bandwidth, 32 GB of device memory, hardware-accelerated shared-memory atomics, and kernel launches cheap enough to run several wide kernels per PIC step without measurable gaps. For 2D electron-scale problems

the *entire field state* now fits in cache several times over, and 10^8 – 10^9 particles fit in device memory. This opens room for a different design philosophy:

a specialized, GPU-native kinetic simulation framework that optimizes time-to-science on one GPU — high single-device throughput, plus an architecture in which new physics capabilities are added as runtime modules rather than code forks.

ArcWarden is that framework. It makes three claims, which organize this paper:


1. **Performance** (Section 5). On identical physics with matched diagnostics, ArcWarden completes a 156-million-particle, 207,000-step kinetic run $2.6\times$ faster than OSIRIS-CUDA on the same GPU; an ablation chain attributes the gain to each design element, and direct measurements show the classic chunk-pipeline design is overhead-bound, not compute-bound, on cache-rich hardware.
2. **Extensibility** (Section 6). One tested GPU particle engine hosts three field models (electrostatic, spectral Darwin, full-Maxwell Yee), full- f and δf kinetic species, a linearized cold-fluid species, prescribed inhomogeneous background fields, loss-cone and anisotropic loading, and absorbing particle and field boundaries — all selected at runtime by plain-text input decks.
3. **Physics capability** (Section 7). The framework passes a three-level validation ladder: text-book verification; reproduction of the published nonlinear whistler-pump trapping simulations of An et al. [8] from their table of parameters; and self-consistent generation of coherent whistler waves with rising-tone frequency chirping in an inhomogeneous magnetic field, following the setup of Tao et al. [9].

ArcWarden does not aim to replace massively parallel general-purpose PIC frameworks; Section 8 states the division of labor explicitly. The niche it targets — electron-scale kinetic studies that fit one modern GPU, run in minutes to hours, and change physics configuration daily — is, we argue, both large and underserved.

2 ArcWarden architecture

2.1 Overall design

ArcWarden is a C++20/CUDA code for a single GPU, organized as a header-only library templated on a compile-time configuration (spatial and velocity dimensionality, particle shape, precision, deposit policy) with thin translation units for kernels. A simulation is defined entirely by a runtime *input deck* (INI text): grid, time step, field model, species list, background field, drivers, boundaries, and diagnostics. Adding an experiment is a new deck, never a new `main()`. The layering is:



[replace with a drawn architecture figure (Fig. 1): deck $\rightarrow$ configuration $\rightarrow$ core $\rightarrow$ kernels, with the physics-module column of Table 4 hanging off the configuration layer]

Fields and moments live on a 2D grid; particle momenta keep all three components (2D3V). All physics uses $\omega_{pe} = 1$, $m_e = 1$, $\varepsilon_0 = 1$ normalization (the 1D chirping module uses $\Omega_{e0} = c = 1$, the convention of the chorus literature).

2.2 Particle representation and memory layout

Particles are a structure of arrays — positions, momenta $\mathbf{u} = \gamma\mathbf{v}/c$, cell index, and (in the δf branch) marker weights — in 32 bit floating point, with all reductions and diagnostics accumulated in 64 bit. Every owning container exposes a POD *view* (`DeviceView`: pointer + size) that is passed to kernels by value, keeping kernels free of host types and ownership. Loading supports quiet starts (coherent sub-cell velocity bases, flat initial charge density) and noisy starts from a counter-based RNG, per species.

The central memory-layout decision is that **fields live in L2 cache**. The full 2D field state — six Yee components plus three current components at 625^2 in single precision — is ≈ 14 MB, 15% of the 96 MB L2 on the benchmark device. Random-order particle gathers are L2-served after first touch, so ArcWarden performs *no* shared-memory staging of field tiles — the core trick of classic GPU-PIC. A microbenchmark shows staging is in fact a net loss on this hardware: a shared-memory-staged push runs at $0.82\times$ the plain L2-served push.

2.3 Fused gather–push–deposit kernel

The entire particle update is one kernel: (i) staggered linear interpolation of E and B from the grid (global loads, L2-served) plus the analytic background B_0 ; (ii) Boris rotation and position update; (iii) Esirkepov current scatter over the 4×4 union stencil; (iv) migration fused into write-back — periodic wrap and cell-index recompute as the position is stored, exactly once. A full PIC step is 4–5 wide kernels with no host round-trips; at $\sim 3\mu\text{s}$ launch cost against $\sim 10\text{ms}$ of work per step, launch gaps are unmeasurable, so no megakernel or persistent kernel is needed.

2.4 Current deposition: shared-memory tile aprons

The scatter is the classic PIC bottleneck: with global atomics each CIC/Esirkepov particle issues ~ 20 atomic adds to scattered addresses. ArcWarden tiles the grid into 16×16 -cell tiles; each tile’s thread block accumulates current in a shared-memory apron (tile + drift margin $D = 2 + \text{stencil halo}$), then flushes it to global memory once per block ($\sim 2\text{k}$ cells) — a cost independent of the number of particles in the tile. The scatter core is one templated routine, `esirkepov_scatter<Sink>`; the `Sink` policy selects shared-tile or global-atomic accumulation, so the fast path and the fallback share a single tested implementation:

Listing 1: Structure of the fused particle kernel (one block per 16×16 -cell tile).

```
__global__ k_push_esirkepov_tiled(particles, bins, fields, ...):
__shared__ float s_jx[...], s_jy[...], s_jz[...]; // tile + apron
zero shared tile; __syncthreads();
for each particle t of this tile (grid-stride over the bin):
    (E,B) = staggered_gather(fields, x0, y0); // global loads -> L2
    u = boris(u, E, B + B0(x0,y0)); // push + background
    (x1,y1) = (x0,y0) + v*dt;
    if stencil(x0..x1) inside apron:
        esirkepov_scatter<SharedJSink>(...); // fast path
    else:
```

```

    esirkepov_scatter<GlobalJSink>(...); // stale-sort stray
    store wrapped positions + recomputed cell index; // fused migration
    __syncthreads();
    flush shared apron to global J (one atomicAdd per apron cell);

```

2.5 Particle sorting: amortized and stale-tolerant

Tile locality comes from a physical counting sort (a reorder of all particle arrays keyed on cell index) that runs every $N_{\text{sort}} \sim 20$ steps — not per step. Two facts make this safe and fast:

- **Correctness does not depend on sort recency.** A particle that has drifted beyond its tile apron since the last sort is detected at scatter time and deposits through the global-atomic sink. The charge-conservation test gates exactly this path: with a deliberately 4-step-stale sort and active strays, the discrete continuity residual is 4.6×10^{-6} .
- **The sort amortizes to noise.** The full sort costs 12.4 ms; at $N_{\text{sort}} = 20$ that is 0.6 ms per ~ 10 ms step, and the end-to-end rate is flat for $N_{\text{sort}} = 10\text{--}40$ ($1.45\text{--}1.54 \times 10^{10}$ particle-steps/s) — there is no cliff to tune against.

This replaces the entire per-step find-movers/exchange/compact pipeline of chunk-based designs with one amortized kernel plus a branch in the scatter; the sort cadence is purely a performance dial.

2.6 Field solvers

Three field branches share the particle engine, selected at runtime by the deck (`[field] model = es | darwin | yee`): a spectral electrostatic solver, a spectral Darwin (magnetoinductive) solver, and a full-Maxwell FDTD (Yee) solver with exactly charge-conserving Esirkepov deposition. Their numerics are given in Section 3; their roles are complementary — the Darwin branch removes the light-wave CFL limit and is the fastest route to low-frequency electromagnetic physics, while the Yee branch provides full-Maxwell ground truth. The two are cross-validated against each other (Section 4.1). All spectral solves are cuFFT-diagonal; the entire field solve at 625^2 costs < 0.2 ms per step, negligible against the particle step.

2.7 Diagnostics and I/O

Diagnostics are deck-selected modules sharing the engine’s data views: mode power and $\delta B/B_0$ spectra, k – t and ω – k reductions, phase-space frames and movies, field snapshots, and energy/charge histories (accumulated in double precision). Heavy reductions run on-device; only reduced arrays cross to the host. Checkpointing serializes the full particle and field state with the run metadata needed for exact restart.

2.8 Correctness gates

Every performance path in the engine is admitted only behind a `cctest` gate: *J-identity* (one tiled step reproduces the flat kernel’s current field to 2×10^{-7} , worst cell); *continuity* ($\partial_t \rho + \nabla \cdot J = 0$ discretely to 4.6×10^{-6} with a stale sort and active strays); *energy parity* (tiled-path field-energy history matches the flat path to 4×10^{-7} through linear growth); and *physics reproduction* (the benchmark figures of Sections 5–7 act as end-to-end acceptance tests). The test suite runs 30+ gates from FFT round-trip up to published-figure reproduction.

3 Numerical models

3.1 Particle equations and interpolation

Particles obey the relativistic Lorentz equation, integrated with the Boris rotation [10] in leapfrog arrangement; velocities are rolled back half a step at initialization. Interpolation is linear (CIC) on gather and scatter in all results shown; the deposit framework is templated over the assignment function, so quadratic (TSC) shapes drop in with a one-cell-wider apron.

3.2 Electrostatic and Darwin spectral solvers

On the periodic domain the electrostatic branch solves $\nabla^2\phi = -\rho/\varepsilon_0$, $E_L = -\nabla\phi$ in k -space. The Darwin branch retains the full magnetoinductive field: with $E = E_L + E_T$, $\nabla \times E_L = 0$, $\nabla \cdot E_T = 0$, the Darwin approximation drops the transverse displacement current, giving the elliptic system

$$\nabla^2\phi = -\rho/\varepsilon_0, \quad E_L = -\nabla\phi, \quad (1)$$

$$\nabla^2 B = -\mu_0 \nabla \times J, \quad (2)$$

$$\nabla^2 E_T = \mu_0 \partial_t J_T, \quad (3)$$

with (3) closed by the acceleration ($\dot{J}$) and momentum-flux ($\langle \mathbf{u}\mathbf{u} \rangle$) moments in the standard Darwin formulation [11], solved in one shot in k -space via the resummed transverse Green's function. No light wave exists in the model, so the time step is set by $\omega_{pe}\Delta t \lesssim 0.2$ and particle transit rather than the light CFL condition — a $6.9\times$ larger step than the Yee branch on the whistler problems of Section 7. This model lineage (UPIC) is the one used by An et al. [8], which makes their published simulations a like-for-like validation target.

3.3 Full-Maxwell FDTD solver

The Yee branch [12] advances all six field components with the full Maxwell equations; current is deposited with the Esirkepov density-decomposition scheme [13], which satisfies the discrete continuity equation to round-off, so no Poisson clean is ever required ($\nabla \cdot B$ stays at round-off; the Gauss residual stays frozen at its initial value, $\sim 2 \times 10^{-7}$). An optional binomial current filter ($[1, 2, 1]^{\otimes 2}/16$ per pass, matching the OSIRIS `smooth` block) suppresses grid-scale noise; being linear and shift-invariant it preserves the discrete Gauss law. A uniform or prescribed background field B_0 enters at the gather, so the Yee update carries only self-consistent fields.

3.4 One-dimensional electron-hybrid model

For field-aligned whistler studies the framework provides a 1D electron-hybrid branch in the Katoh–Omura model class [14, 2]: transverse fields on a staggered 1D grid, a linearized cold electron fluid carrying the whistler dispersion, relativistic kinetic hot electrons (full- f or nonlinear δf), the adiabatic mirror force of a prescribed inhomogeneous field $B_0(h) = B_0(1 + ah^2)$ (or a dipole profile), particle reflection at the domain ends, Umeda-style masked absorbing layers for the fields [15], and an equatorial triggering antenna. This is the branch used for the chirping demonstration (Section 7.2).

3.5 δf representation

Hot species may be represented as δf markers about an analytic equilibrium f_0 (bi-Maxwellian, optionally loss-cone-subtracted, with an (E, μ) -mapped mirror equilibrium along inhomogeneous

B_0); marker weights evolve along characteristics and deposit $\delta\rho, \delta J$. The δf branch is gated by equilibrium-hold tests (weights stay at machine noise in a held equilibrium) and by growth-rate agreement with full- f on the same deck.

4 Verification

The first level of the validation ladder asks only: *is the code correct?* The `cctest` suite covers, per branch:

- *Electrostatic*: cold Langmuir oscillation frequency; two-stream instability growth rate against the fluid dispersion; bump-on-tail relaxation; Landau damping (δf).
- *Darwin*: magnetostatic B -from- J against the analytic current sheet; Weibel instability growth rate.
- *Yee*: vacuum wave propagation; Langmuir oscillation on the FDTD path; cold-plasma whistler eigenmode frequency (0.03% of the full-Maxwell root); magnetized single-particle gyromotion and mirror bounce; long-run energy conservation; discrete continuity and Gauss residuals (Section 2.8).
- *Hybrid/chirping*: cold whistler dispersion on the hybrid grid; anisotropy-driven growth rate against the kinetic dispersion solver; δf equilibrium hold under inhomogeneous B_0 .

Figure 1 condenses the electrostatic-branch level of the ladder into one deck-driven figure (produced by `scripts/repro/plot_verification.py`): the two-stream instability grows at $\gamma = 0.343\omega_{pe}$ against the cold-beam rate $\omega_{pe}/\sqrt{8} = 0.354$ (the 3% deficit is the finite $v_{th}/v_0 = 0.1$ thermal correction), saturates into the classic trapping vortex, and over a $10\times$ longer run the total energy oscillates within $\pm 2 \times 10^{-3}$ with no secular drift while the net charge is conserved to round-off. [extend to a companion EM panel set (whistler eigenmode + Weibel growth) from the same script]

4.1 Cross-branch validation

The Darwin and Yee branches are validated against each other where the Darwin approximation holds ($\omega^2 \ll k^2 c^2$); agreement between a radiation-free spectral solver and a full-Maxwell FDTD solver sharing only the particle engine is a sensitive check on both. At eigenmode level, a cold parallel whistler at $kc = 3\omega_{pe}$, $\omega_{ce} = 0.5\omega_{pe}$ gives $\omega = 0.44980$ (Darwin; analytic 0.45000) and 0.44885 (Yee; full-Maxwell root 0.44897) — and the 0.21% gap *between* the measured frequencies matches the 0.23% analytic displacement-current gap between the two field models: the branches differ by exactly the physics that separates the models. At nonlinear level, the An et al. Sim 3 deck was run through both branches (the Yee branch at its light-CFL step, $6.25\times$ more steps): whistler amplitude peaks at $\delta B/B_0 = 0.105$ (Darwin) vs 0.108 (Yee) with identical ring-down, and the nonlinear electrostatic band and trapping structures agree (Fig. 2).

5 GPU performance

All measurements are on one NVIDIA RTX 5090 (Blackwell, `sm_120`: 170 SMs, 96 MB L2, 31.4 GB GDDR7), CUDA 13.3. The benchmark problem is case 7 of the whistler/electron-acoustic parameter study of Ma et al. [16]: a bi-Maxwellian electron plasma, $T_\perp/T_\parallel = 5$, $\beta_\parallel = 0.0091$,




Figure 1: Electrostatic-branch verification from one input deck (decks/verification/two_stream.ini): (a) field-energy growth vs the analytic two-stream rate; (b) phase space at saturation; (c) relative total-energy drift and (d) relative net-charge drift over a 6000-step run.




Figure 2: Darwin (spectral) vs Yee (full-Maxwell) branch on the same input deck (An et al. Sim 3): whistler amplitude history (left), secondary electrostatic band power over the noise floor (middle), and Yee-branch phase space at $t = 1200 \omega_{pe}^{-1}$ (right).

$\omega_{ce}/\omega_{pe} = 0.25$, 625^2 cells, 400 particles per cell (156.25M particles), run to $t = 3000 \omega_{pe}^{-1}$ (206,897 Yee steps).

5.1 Head-to-head against OSIRIS-CUDA

Both ArcWarden (Yee branch, linear CIC + 3-pass binomial filter) and OSIRIS 4.4.4 CUDA (quadratic + binomial-3, its production configuration) ran the case end-to-end with matched diagnostics and were reduced by the same eight-panel analysis (Fig. 3). The physics agrees panel by panel: whistler onset at $k_x \simeq 2.8 \omega_{pe}/c$ in both, wave-normal angle 44–47° in both, the electron-acoustic band above noise by the same factor in both, the beam-mode ridge on $\omega = 0.0546c k_x$ in both; field-energy histories agree to four significant figures through the linear phase.

Table 1: End-to-end wall time, identical physics and diagnostics. 156.25M particles, 206,897 steps, same RTX 5090.

code	wall time	particle-steps/s	speedup
OSIRIS 4.4.4 CUDA (quadratic, binomial-3)	5392.8 s	6.0×10^9	1
ArcWarden Yee tiled (linear, binomial-3)	2079.7 s	1.55×10^{10}	2.59×

The interpolation orders differ (quadratic is OSIRIS’s production configuration; ArcWarden is linear+filter), but Section 5.3 shows by direct measurement that interpolation order moves the OSIRIS-CUDA rate by only 7%, so the comparison is dominated by architecture, not arithmetic. [optional: quadratic shapes in the tiled framework (est. 20–30% cost) for exact shape parity]

5.2 Ablation: where the speedup comes from

Table 2: Ablation chain, benchmark configuration, diagnostics off.

configuration	particle-steps/s	gain
flat: fused kernel, global-atomic Esirkepov	4.5×10^9	—
+ 16×16 tile sort (every 20) + shared apron deposit	1.51×10^{10}	3.4×
+ migration fused into scatter write-back	1.61×10^{10}	+6.6%

Two supporting microbenchmarks: the tiled charge deposit alone is 6.4× faster than global atoms (14.9 vs 2.3 Gdep/s); the plain L2-served push runs at 35.4 Gpush/s while the shared-memory field-staging variant of the same push measures 0.82× — the negative result behind the fields-live-in-L2 rule (Section 2.2). Sort cadence is a flat dial ($N_{\text{sort}} = 10/20/40 \rightarrow 1.45/1.51/1.54 \times 10^{10}$). [clean single-session ablation rerun, same binary, same night]

5.3 The classic design is overhead-bound on this hardware

Three independent measurements on the OSIRIS-CUDA reference identify per-step particle book-keeping, not computation, as its pacing item: (i) GPU utilization is 52% during the production run, sampled over minutes; (ii) interpolation order is nearly free — switching linear→quadratic ($\sim 2\times$ deposit arithmetic, $2.25\times$ stencil) moves the rate only from 6.3 to 5.9×10^9 (−7%), which a compute-bound code cannot do; (iii) chunk size is an exact null — doubling `chunk_size` 512→1024 leaves the rate at 5.92×10^9 in both cases, so the cost scales with the pipeline stages themselves, not




Figure 3: The benchmark case run in ArcWarden and OSIRIS-CUDA with matched diagnostics; identical eight-panel reduction. [final layout: possibly 4 selected panels each]

chunk granularity. The per-step find/exchange/compact pipeline is precisely the term ArcWarden deletes (Section 2.5). During the same benchmark ArcWarden sustains 100% GPU utilization at ≈ 575 W.

5.4 Kernel-level budget

The step decomposes as $T_{\text{step}} = T_{\text{sort}} + T_{\text{deposit+push}} + T_{\text{field}} + T_{\text{diag}}$. On the Yee branch the fused particle kernel is 87.7% of all GPU time (median 9.3 ms/step); the amortized sort contributes 0.64 ms/step; the entire Maxwell update including nine filter passes totals ≈ 0.05 ms/step (0.5%). The GPU spends essentially all of its time on the irreducible particle work. Table 3 gives the equivalent budget for the Darwin branch, whose per-step tile sort is a documented legacy (36% of the step) with a validated upgrade path to the amortized scheme. [\[ncu report: SOL memory/compute %, occupancy, L2 hit rate of the gathers \(profiling/collect_ncu.sh, needs perf-counter permission\)\]](#)

Table 3: Darwin-branch per-step kernel budget at benchmark scale (Nsight Systems averages over 300 steps).

kernel	ms/step	share
per-step tile sort (legacy; to be amortized)	11.6	36%
fused gather + Boris push	8.0	25%
$\rho+J$ tiled deposit	3.4	11%
$\langle uu \rangle$ (amu) moment deposit	3.3	10%
$\dot{J}$ (dcu) moment deposit	3.0	9%
migration (legacy; fused into scatter in the Yee branch)	2.1	7%
spectral solves (all cuFFT + k -space kernels)	0.17	0.5%

5.5 Model choice as a performance multiplier: the Darwin branch

The Darwin branch runs the benchmark-scale problem at $\Delta t = 0.1 \omega_{pe}^{-1}$ — $6.9\times$ the Yee light-CFL step — at a measured 5.0×10^9 particle-steps/s, so the full benchmark physical time costs ≈ 943 s: per unit of *physical time simulated*, $2.2\times$ cheaper than the already accelerated Yee branch and $5.7\times$ cheaper than the OSIRIS-CUDA reference, on problems whose physics the model captures. The point is architectural: on a single GPU, choosing the minimal field model for the physics at hand *stacks* with every kernel optimization, and the modular solver layer (Section 2.6) is what makes that choice a one-line deck edit rather than a different code. [\[scaling curves: rate vs ppc \(100–4096\) and vs grid \(256²–2048²\), both codes — planned Fig.\]](#)

6 Modular physics capabilities

The second claim of the paper is that the performance of Section 5 does not come at the price of rigidity. All capabilities in Table 4 share the same particle engine, deposit framework, and correctness gates; each was added as a module against the tested core, and each is selected at runtime by a deck. A single physics setup can be moved between field models by editing one line (`[field] model = es | darwin | yee`), which is also how the cross-branch validations of Section 4.1 are run.

Two design choices make the module set cheap to extend. First, the compile-time configuration (SimConfig) specializes the kernels — precision, shape, deposit policy, background-field capability

Table 4: Physics modules and where each is demonstrated.

capability	implementation	demonstrated application
electrostatic	spectral Poisson	two-stream, bump-on-tail (§4)
Darwin EM	spectral, moment-closed (3)	whistler pump / An et al. (§7.1)
full EM	Yee + Esirkepov	anisotropy whistlers + EAW (§5.1)
inhomogeneous B_0	prescribed $B_0(h)$: parabolic, dipole	mirror-trapped electrons (§7.2)
anisotropic / loss-cone f_0	configurable loader; (E, μ) mirror map	chorus excitation (§7.2)
δf species	weight evolution about analytic f_0	low-noise growth, seeded runs
hybrid kinetic-fluid	linearized cold fluid + hot PIC	1D chorus model (§3.4)
absorbing field boundary	Umeda masked layers [15]	open field lines (§7.2)
particle boundary	reflection / absorption at domain ends	bounce motion, precipitation
wave driver	traveling-wave pump; triggering antenna	An et al.; triggered emissions

— so a module pays only for what it uses; a run without B_0 compiles to the same kernel it had before B_0 existed. Second, the correctness gates of Section 2.8 are cumulative: every module lands with its own gate (e.g. the δf mirror equilibrium hold, the boundary-reflection benchmark with $R < 1\%$ across the whistler band), and all previous gates must keep passing, so the physics surface grows without eroding the tested core.

7 Advanced nonlinear demonstrations

The remaining two levels of the validation ladder move from *correct* to *scientifically capable*: reproduction of an established nonlinear kinetic benchmark, then a demonstration at the frontier of what kinetic codes are asked to do.

7.1 Published nonlinear benchmark: whistler-driven electron trapping (An et al. 2019)

The Darwin branch reproduces the whistler-pump simulations of An et al. [8] — a like-for-like model comparison, since that work used the spectral-Darwin UPIC code, the same field model ArcWarden implements. All three simulations of their Table I run from decks carrying the published parameters verbatim: a traveling-wave pump drives a quasi-parallel whistler; when its amplitude crosses $\delta B/B_0 \approx 0.1$, cyclotron-resonant trapping deforms the electron distribution and a secondary electrostatic band erupts — beam-mode Langmuir waves when the resonant velocity sits far out on the tail ($v_r = 3.2 v_{th}$, Sim 1), electron-acoustic waves and unipolar structures at mid-distribution (Sim 2), and phase-space holes with bipolar fields when it sits in the thermal core ($v_r \approx v_{th}$, Sim 3). ArcWarden reproduces the mechanism and the amplitude threshold in all three cases; Fig. 4 shows Sim 3 at 268M particles: the chain of trapping vortices at $v_r = 1.04 v_{th}$ with the bipolar $\delta E_{\parallel}$ signature aligned to the holes, the electrostatic band erupting at $28\times$ the noise floor once the whistler crosses threshold. One honest caveat: matching the published $\delta B/B_0$ trajectory requires scaling the pump amplitude $\sim 5\text{--}6\times$ over the Table I values (An et al. likewise tuned the drive); the physics downstream of the wave amplitude is parameter-free.

This benchmark certifies more than a growth rate: phase-space structure, nonlinear trapping, electron holes, and broadband secondary spectra — the nonlinear vocabulary the framework exists to compute.




Figure 4: An et al. [8] Sim 3 on the Darwin branch (268M particles): bipolar $\delta E_{||}(x)$ (top), electron phase space $f(x, v_{||})$ (middle), and signed δf (bottom) showing the chain of trapping vortices at $v_r = 1.04 v_{th}$.

7.2 Capability demonstration: self-consistent coherent whistler waves with frequency chirping in an inhomogeneous plasma

As the top of the ladder we demonstrate a composite capability that exercises most of Table 4 at once: an inhomogeneous (mirror) background field $B_0(h)$, an anisotropic hot electron population held in a mirror equilibrium, absorbing field boundaries and reflecting particle boundaries, and a triggering antenna — the setup of Tao et al. [9] for rising-tone chorus. From these published parameters the 1D electron-hybrid branch self-consistently generates coherent whistler-mode wave elements whose frequency *chirps* upward, together with the resonant phase-space hole whose motion accompanies the sweep (Fig. 5); the element timing, chirp rate, and subpacket structure match the published run, and a δf null test (no anisotropy drive $\rightarrow$ no element) confirms the emission is physical, not numerical.

We present this as a *capability demonstration*, not a study of the chirping mechanism: frequency chirping is exquisitely sensitive to errors in resonant particle dynamics [14, 17, 2], which makes it the strongest integration test available for a whistler code — every module in the chain (mirror force, equilibrium load, boundaries, cold-fluid dispersion, hot-particle resonance) must be right simultaneously, or the tone does not form. Physics questions about chorus generation are deliberately out of scope here and are the subject of separate work.

8 Discussion

Design philosophy and division of labor. ArcWarden does not aim to replace massively parallel general-purpose PIC frameworks; OSIRIS, VPIC, PIConGPU, and WarpX serve problems that need thousands of GPUs, and nothing here competes with that. ArcWarden optimizes a different quantity: *time-to-science on one GPU* — the wall-clock and iteration cost of an electron-scale kinetic study that fits a single modern device, including the cost of adding the next physics module. The benchmark result should be read in that frame: a 156M-particle, 207k-step kinetic run finishing in 35 minutes on a \$2k consumer device changes what a single researcher can scan in a week.

Scope of the L2 argument. “Fields live in L2” holds for 2D electron-scale grids up to roughly 1600^2 – 2000^2 on a 96 MB device. For larger grids or 3D, the gather working set exceeds L2 and tile-locality of field reads becomes valuable again — but via the same amortized stale-tolerant sort already present, not via per-step chunk bookkeeping; the stale-tolerant-sort result is the part we expect to transfer unchanged to any architecture, including the exascale codes.

Comparison fairness. OSIRIS ran its production configuration (quadratic shapes); ArcWarden ran linear+filter. The overhead-bound measurements (Section 5.3) bound the effect of this difference at $\sim 7\%$ of the OSIRIS rate; a quadratic tiled deposit in the Sink framework is the cleanest way to close the gap entirely. [\[decide before submission\]](#)

Limitations and roadmap. All results are single-GPU, single-precision hot path, 2D3V (1D for the hybrid branch). The Darwin branch still sorts every step (36% of its budget, Table 3) with a validated upgrade path. Planned extensions — multi-GPU domain decomposition and the dipole-field L-shell geometry — reintroduce halo exchange, which the fused kernel tolerates by design (the stray fallback already handles out-of-tile particles).




Figure 5: Electron-hybrid branch, setup of Tao et al. [9]. Top: dynamic spectrum $B_w(f, t)$ of the self-consistently generated rising-tone element. Bottom: tracking of the resonant phase-space hole that accompanies the frequency sweep. [final figure selection/cropping; consider adding the $B_0(h)$ profile + loss-cone f_0 as a setup panel]

Development methodology. ArcWarden was developed by one physicist working with an LLM coding agent (Claude Code, Anthropic) under explicit human-written discipline: milestones with quantitative acceptance gates stated before the work, correctness tests preceding any optimization, a committed performance baseline every kernel change must beat, and published figures as acceptance tests. The gates of Section 2.8 and the validation ladder of this paper are that discipline made visible; the full development history is preserved in the repository. We found the human role concentrated into verification — writing gates, judging physics plausibility, catching stability-limit violations — and we note the practice as one reproducible answer to the question of how agent-written GPU code can be trusted in scientific computing.

9 Conclusions

ArcWarden demonstrates that a PIC framework designed around what modern GPUs actually provide — cache-resident fields, fused step kernels, shared-memory scatter aprons, and amortized stale-tolerant sorting — is simultaneously fast, extensible, and scientifically capable: 2.6× the throughput of an established GPU-PIC framework on identical physics, with the gain accounted element by element; a runtime-modular physics layer in which field models, background geometries, distributions, and boundaries share one tested particle engine; and a validation ladder that runs from textbook dispersion through published nonlinear benchmarks to self-consistent coherent whistler chirping in an inhomogeneous plasma. Electron-scale kinetic studies that were cluster jobs are now interactive desktop work, and the framework is built so the next physics module lands as a deck and a gate, not a fork.

Acknowledgments

Development used Claude Code (Anthropic) under the methodology of Section 8. [grants, computing, colleagues]

Code and data availability

ArcWarden is available at [repository URL on publication]. All input decks (grouped per paper section in `decks/`), analysis scripts, profiling reports, and the OSIRIS input files used for the comparison are included in the repository; every figure in this paper maps to a named deck (see `decks/README.md`).

References

- [1] C. K. Birdsall and A. B. Langdon. *Plasma Physics via Computer Simulation*. CRC Press, 2004.
- [2] Y. Omura. Nonlinear wave growth theory of whistler-mode chorus and hiss emissions in the magnetosphere. *Earth, Planets and Space*, 73:95, 2021.
- [3] R. A. Fonseca, L. O. Silva, F. S. Tsung, V. K. Decyk, W. Lu, C. Ren, W. B. Mori, S. Deng, S. Lee, T. Katsouleas, and J. C. Adam. OSIRIS: a three-dimensional, fully relativistic particle in cell code for modeling plasma based accelerators. In *Computational Science — ICCS 2002*, volume 2331 of *Lecture Notes in Computer Science*, pages 342–351. Springer, 2002.

- [4] K. J. Bowers, B. J. Albright, L. Yin, B. Bergen, and T. J. T. Kwan. Ultrahigh performance three-dimensional electromagnetic relativistic kinetic plasma simulation. *Physics of Plasmas*, 15(5):055703, 2008.
- [5] M. Bussmann, H. Burau, T. E. Cowan, A. Debus, A. Huebl, G. Juckeland, T. Kluge, W. E. Nagel, R. Pausch, F. Schmitt, U. Schramm, J. Schuchart, and R. Widera. Radiative signatures of the relativistic Kelvin–Helmholtz instability. *Proceedings of SC13: International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 5:1–5:12, 2013. (PConGPU).
- [6] L. Fedeli, A. Huebl, F. Boillod-Cerneux, T. Clark, K. Gott, C. Hillairet, S. Jaure, A. Leblanc, R. Lehe, A. Myers, C. Piechurski, M. Sato, N. Zaïm, W. Zhang, J.-L. Vay, and H. Vincenti. Pushing the frontier in the design of laser-based electron accelerators with groundbreaking mesh-refined particle-in-cell simulations on exascale-class supercomputers. In *SC22: International Conference for High Performance Computing, Networking, Storage and Analysis*, 2022. (WarpX; ACM Gordon Bell Prize).
- [7] V. K. Decyk and T. V. Singh. Particle-in-cell algorithms for emerging computer architectures. *Computer Physics Communications*, 185(3):708–719, 2014.
- [8] X. An, J. Bortnik, B. Van Compernelle, V. Decyk, and R. Thorne. Unified view of nonlinear wave structures associated with whistler-mode chorus. *Physical Review Letters*, 122:045101, 2019.
- [9] X. Tao, F. Zonca, and L. Chen. Identify the nonlinear wave-particle interaction regime in rising tone chorus generation. *Geophysical Research Letters*, 44:3441–3446, 2017.
- [10] J. P. Boris. Relativistic plasma simulation — optimization of a hybrid code. In *Proceedings of the Fourth Conference on Numerical Simulation of Plasmas*, pages 3–67. Naval Research Laboratory, Washington, DC, 1970.
- [11] C. W. Nielson and H. R. Lewis. Particle-code models in the nonradiative limit. *Methods in Computational Physics*, 16:367–388, 1976.
- [12] K. S. Yee. Numerical solution of initial boundary value problems involving Maxwell’s equations in isotropic media. *IEEE Transactions on Antennas and Propagation*, 14(3):302–307, 1966.
- [13] T. Zh. Esirkepov. Exact charge conservation scheme for particle-in-cell simulation with an arbitrary form-factor. *Computer Physics Communications*, 135(2):144–153, 2001.
- [14] Y. Katoh and Y. Omura. Computer simulation of chorus wave generation in the Earth’s inner magnetosphere. *Geophysical Research Letters*, 34:L03102, 2007.
- [15] T. Umeda, Y. Omura, and H. Matsumoto. An improved masking method for absorbing boundaries in electromagnetic particle simulations. *Computer Physics Communications*, 137:286–299, 2001.
- [16] D. Ma, X. An, X. Tao, J. Bortnik, X.-J. Zhang, and V. Angelopoulos. Nonlinear interactions between electron acoustic waves and electron cyclotron harmonic waves / whistler waves. *Physics of Plasmas*, 31:022304, 2024. arXiv:2308.03938; TODO verify final title.

- [17] M. Hikishima, S. Yagitani, Y. Omura, and I. Nagano. Full particle simulation of whistler-mode rising chorus emissions in the magnetosphere. *Journal of Geophysical Research*, 114:A01203, 2009.